

IssueOps

Continuous AI for GitHub Issues.

"We already burn countless unpaid hours triaging issues and dealing with duplicates instead of writing code."

— Mike McQuaid, Lead Maintainer of Homebrew

Team: VaaS



DS 252 : Introduction to Cloud Computing

September 17, 2025

The Problem

Developer productivity is being drained by GitHub Issues.

7.3 hours lost per week
per developer to triage, duplicates, and notifications

3+ day delays
before issues are even categorized or routed

59% of maintainers
consider quitting due to burnout and overhead

Market Opportunity

- 150M+ developers on GitHub
- 420M+ repositories hosted
- \$37.45B automation market by 2030
- \$21.6M TAM (GitHub-specific)
- 9.52% CAGR growth rate

Revenue Potential

- Team Subscriptions: **\$14.4M/year**
- Individual Freemium: **\$6.0M/year**
- Enterprise Plans: **\$1.2M/year**
- **Total TAM: \$21.6M annually**

Our Solution

IssueOps Agents automate the GitHub issue lifecycle end-to-end:
Intake → Triage → Assignment → Notification.

Our Core Ideas

- **Connected Workflow:** track each issue as it moves across agents; preserve context and state.
- **Smart Triage:** combine AI understanding with lightweight rules for consistent decisions.
- **Elastic Scaling:** containerized services that scale with repo activity; no manual ops burden.

Key Value Proposition

- Cut triage delays from days to hours; surface priorities fast.
- Reduce notification noise with batching, rate limits, and targeted alerts.
- Clear ownership and next actions through labels, assignees, and linked duplicates.

How it Works

End-to-End Flow


- When a user installs the app we do a one time scan of the repo and send the file to our server, which stores context for the future.
- When a new issue is created on GitHub, a webhook sends the issue details to our service.
- Our service receives the data, processes it, and determines the correct label using AI.
- The label is applied back to the issue on GitHub through the API.

Data Movement

- **From:** GitHub (issue title, description, metadata).
- **To:** Our server for analysis and decision-making (might send data as lambda functions to reduce load).
- **Back to:** GitHub with updated labels and status.

GitHub Issue \Rightarrow Webhook \Rightarrow Our Service \Rightarrow GitHub (Updated)

Intake Agent




The Intake Agent is the first gatekeeper. It cleans, organizes, and enriches every incoming issue so only actionable work reaches developers.

Key functions

1. **Spam Shield:** Blocks bots, junk, and irrelevant noise before it clutters the repo.
2. **Duplicate Finder:** Detects similar issues and links them into one clear thread.
3. **Stale Detector:** Flags outdated or abandoned issues, keeping the backlog lean.
4. **TODO Tracker:** Scans code for TODO notes and turns them into GitHub issues automatically.
5. **Repo Validator:** Ensures bug reports include clear steps for reliable reproduction.

Triaging Agent



The Triage Agent first enriches incoming issues with relevant metadata, to build a complete picture. This enriched context allows the agent to accurately classify the issue and set a precise priority

Key functions


1. **Metadata Enrichment:** Adds context like environment, dependencies, severity, and history to speed up fixes.

```
// Input from triage agent
{
  heading = "Triage alert";
  comment = "Triage handles user care alerts",
  code = "Effects == triage //|| NoOp == null",
  logpointer
}

// Enriched output
{
  heading = "Authentication Failure in Logon Monitor";
  comment = "Logon monitor handles cases where credentials are valid,
             but user can't be found or password is incorrect";
  description: "Module is in charge of RDP-NL
               and RDP-UI authentication";
  code = "if (auth == null) {
    throw new Error('Auth object missing');
  }
  logpointer"
}
```

2. **Classification:** AI + rules map issues to components and labels instantly.
3. **Priority Scoring:** Assigns P0-P3 based on impact and urgency.

Assignment Agent




The Assignment Agent maps triaged issues to the best person or team, respecting skill, availability, and SLA constraints.

Key functions

1. **Skill Match:** Maps labels and components to contributors' expertise so the right person gets the job.
2. **Workload Balance:** Distributes assignments to avoid overload and keep throughput steady.
3. **Availability Check:** Respects calendars, on-call status, and recent activity when choosing assignees.
4. **SLA Escalation:** Automatically escalates or reassigns when timeouts or priority SLAs are breached.

Notification Agent



The Notification Agent delivers alerts with precision and gives teams a live dashboard to track who got which issue, when, and how it's progressing.

Key functions

1. **Smart Channeling:** Routes updates through Slack, Email, or GitHub based on context.
2. **Noise Control:** Batches and rate-limits alerts to prevent fatigue.
3. **Personalized Digests:** Curates updates and escalates urgent items instantly.
4. **Agent Dashboard:** Shows issue ownership, notification history, and response times in real time.

Workflow




Figure 1: Workflow

Tech Stack

Core Platform

- GitHub App (repo-scoped permissions)
- Python + Flask for webhook handling
- Runs in a Docker container (local or cloud)
- Optional: simple cron for retries

Data & Config

- Config in YAML/JSON (labels, owners)
- Optional SQLite/Postgres for logs
- Secrets via .env or GitHub Actions

AI & Logic

- Zero-shot classification via hosted API
- Lightweight rules for mapping predictions
- Confidence threshold for safe automation

Integrations & DevOps

- GitHub REST API for labels/updates
- Optional Slack/Email notifications
- ngrok for local webhook testing
- CI/CD with GitHub Actions
- Basic logs + health check endpoint

Goal: Simple, modular, and easy to deploy on free-tier or university VM.

2-Week MVP: What It Does & What We Are Doing

What the MVP Does

- Takes new issues from GitHub and sends them to our server automatically.
- Runs a simple **4-step classification** on each issue: *duplicate? → issue type (bug/feature/...)* → *assign to the right engineer type* → *notify based on priority/role*.
- Sends the results back to GitHub: labels applied, owners notified; duplicates handled (per repo policy).

What We Are Doing in These 2 Weeks

- **Initial setup:** perform a one-time full scan of the repo and save key data **encrypted** on our server (or a simple free SQL database) for faster, consistent decisions later.
- **Week 1:** connect GitHub to our service (webhook) and receive issue data reliably.
- **Week 2:** implement classification and apply updates back to GitHub (labels, basic duplicate handling, owner assignment + notifications in a minimal form).

GitHub ⇒ Our Server ⇒ 4-Step Classification ⇒ Labels/Assign/Notify ⇒ GitHub (Updated)

Project Timeline and Assignments

Weeks 1-2: Foundational Agent Development

Ashwin: Build Intake (Core). | **Vinay:** Build Triage and Intake (TODO). |
Anmol: Design Assignment. | **Sai Harsh:** Build Intake (Repro Validation).



Weeks 3-4: Full Agent Suite Development

Ashwin: Refine Intake. | **Vinay:** Refine Triage. | **Anmol:** Build Assignment Agent. | **Sai Harsh:** Build Notification Agent.



Weeks 5-6: Full Pipeline Integration

Ashwin: Integrate Intake. | **Vinay:** Integrate Triage. | **Anmol:** Integrate Assignment. | **Sai Harsh:** Integrate Notifications.



Weeks 7-8: Verification and Documentation

Ashwin: Document Intake | **Vinay:** Document Triage | **Anmol:** Document Assignment | **Sai Harsh:** Verification & Document Notifications